

Distilling Contrastive Learning Models

Valerie Lim

vlim31@gatech.edu

Ng Kai Wen

kng71@gatech.edu

Kenneth Lim

klim83@gatech.edu

April 24, 2022

Abstract

Pretrained word embedding Natural Language Processing (NLP) models like BERT are able to provide state of the art results for many NLP tasks, but the area of Semantic Textual Similarity (STS) is a notable exception. Contrastive learning has been shown to help improve the quality of sentence embeddings produced, but these models are often large, since they use BERTbase or RoBERTa as pretrained models. Knowledge distillation meanwhile could address this by making models more lightweight. Hence, to reap both of these benefits, we seek to apply a suitable contrastive learning method based on the SimCSE paper, to a model architecture adapted from a knowledge distillation based model, DistilBERT, in an unsupervised setting. We also engage in extensive hyperparameter tuning, experiment with different pooling layers and automatic mixed precision to boost performance and efficiency. Our final lightweight model DistilFACE achieves an average of 72.1% in Spearman’s correlation on STS tasks, a 34.2% improvement over BERTbase with the default parameters.

1 Introduction

1.1 Problem Statement

Pre-trained word embeddings are a vital part of natural language processing (NLP) systems, providing state of the art results for many NLP tasks as the pre-trained embedding-based models consistently outperform alternatives such as the embedding-layer-based Word2Vec [5]. Accurate word embeddings have huge downstream benefits for tasks related to information retrieval and semantic

similarity comparison [18], especially in a typical industrial setting where labelled data is scarce. As a result, in recent years, word embeddings from pre-trained models such as BERT (Bidirectional Encoder Representations from Transformers) have been gaining traction in many NLP works.

However, studies [6] [18] have shown that BERT does not produce good enough embeddings for semantic textual similarity tasks. As a result, recent breakthroughs [6] [12] have found that a contrastive objective is able to address the anisotropy issue of BERT. At a high level, contrastive learning is a technique where given two augmented samples of x , $x+$ (positive sample), $x-$ (negative sample), we’d like for $x+$ to be close to x and for $x-$ to be far away from x . This improves the quality of BERT sentence representation.

Yet, a problem often associated with BERT and the contrastive loss models derived off it is their large size. There is an increasing need to build machine learning systems that can operate on the edge, instead of calling a cloud API, in order to build privacy respecting systems and prevent having users send sensitive information to servers [1]. Thus, there is a growing impetus for lightweight and responsive models. However, most of the pre-trained models [6] [18] with contrastive loss are still relatively large and computationally heavy, making it hard to deploy on the edge. Deployment on the edge would also allow for reduced time and network latency, which could be critical in certain use cases such as when there is a high QOS (Quality of Service) requirement [1]. A common solution for creating lighter models is via knowledge distillation, where a trained, complex teacher model trains a lighter, shallower student model to generalise [14]. Hence, our goal was to address this

joint problem of lack of lightweight models that can perform well on semantic textual similarity tasks, by creating an unsupervised pretrained model that combines knowledge distillation and contrastive learning to address their respective problems. To do this, we first adopt the architecture and pre-trained checkpoints of a knowledge distillation based model, DistilBERT, to reduce the size of our model. We then adapt from authors who successfully implemented contrastive learning on sentence embeddings via unsupervised learning [6], by feeding samples, termed as x , into the pre-trained encoder twice: by applying standard dropout twice to obtain two different sets of embeddings termed as $x+$. Then we take other sentences in the same mini-batch as $x-$ and the model predicts the positive sample among the negatives. (adapt SimCSE Fig 1a). We subsequently experiment with different pooling layers, automatic mixed precision and hyperparameter tuning to enhance performance.

1.2 Implications

Our group believes there is significant value in experimenting with whether a traditional knowledge distillation method like DistilBERT, can be directly adapted to be combined with an existing contrastive learning method such as SimCSE, without having to create an entirely new model. Such an approach makes such a combination of techniques more accessible. In addition, there are definitely environmental and economical benefits to doing so, given how computational intensive large-scale pretraining of new models are. For instance, 16 TPU chips were required to train BERTbase, with each round of pretraining taking 4 days [4], costing about US\$6,912 each time [16]. Furthermore, success in our project will enable the creation of lightweight pre-trained embedding-based models that work well in semantic textual similarity tasks. This has many applications, such as searching and ranking of relevant documents based on similarity [8], and now it could be even further applied in edge scenarios. For example, such functionality can now be enabled at sensitive government, defence or health institutions, where computers might not be allowed to be connected to the cloud freely [14].

2 Related Work

2.1 Contrastive Learning Models

SimCSE [6] presented a simple contrastive learning framework that vastly improved the state-of-the-art sentence embeddings. By training unsupervised models using contrastive objective, the resulting quality of sentence embeddings improved by 4.2%. The contrastive objective can be understood as maximising the lower bound of mutual information between different views of the data [7]. Ultimately, the goal of contrastive learning is to learn effective representation via pulling semantically close neighbours together and pushing apart non neighbours, through a score function which measures similarity between these two features. This would be elaborated on the solution section. The input sentences are encoded with a pretrained language model such as DistilBERT, which is elaborated upon in the next section.

2.2 Knowledge Distillation Models

In the present work, contrastive learning is applied on a distilled version of BERT, DistilBERT. Sanh et. al [14] was able to reduce BERT model size by 40% while retaining 97% of language understanding capabilities by leveraging on knowledge distillation. DistilBERT is able to retain a significant extent of model performance despite a huge reduction of memory footprint because of the benefits conferred from knowledge distillation. In knowledge distillation, DistilBERT is known as the student, while BERT is known as the teacher. Both have the same architecture. For instance, given this sample input sentence “the <mask> licked its fur and howled”, the teacher model outputs “dog” as the mask, but “wolf” and “fox” also have high scores. The distillation loss between the teacher and the student aims to align the probability distributions, rather than just the hard predictions, which is the output from the teacher model. The probability distribution is a useful addition as the student is able to learn to rank “wolf” and “dog” highly, which is useful info about the world [14].

2.3 Combining Contrastive Learning with Knowledge Distillation

The core key elements of contrastive learning are augmentation techniques to create positive samples, hard negative mining techniques to create negative samples, and sufficiently large batch size [2]. We hypothesise that there is nothing inherent about these properties that would prevent contrastive learning’s effective use when applied together to DistilBERT instead of BERT. We believe that it would be especially meaningful to pursue this combination of contrastive learning and knowledge distillation techniques as these were the two methods that were highlighted as the most dominant and effective surrogate tasks used in unsupervised learning and self supervised learning in the lectures [9].

There has been some work done today to combine the two techniques. The Contrastive Distillation on Intermediate Representations (CoDIR) framework [15] uses contrastive learning to distil the teacher’s hidden states to the student through three objectives, an original training objective, a distillation objective, and an additional contrastive learning objective where the student is trained to distil knowledge through intermediate layers of the teacher. However, this utilisation of both teacher’s output layer and its intermediate layers for student model training is a deviation from traditional knowledge distillation methods, and involves the creation of a new framework. This differs from our intention of checking compatibility of existing methods. This framework is also instead applied directly on RoBERTa, and is more complex to implement. In addition the main objective of applying contrastive learning in CODIR is to compress BERT [15], whereas for us the compression is mainly done through the use of DistilBERT, and application of contrastive learning is mainly for enhancing performance.

2.4 Pooling Methods

The original authors of BERT [4] experimented with different pooling strategies and the top three performers were concat last 4, weighted sum of last four hidden and second-to-last hidden. However, these were on fine-tuned tasks and may not be good benchmarks for semantic quality of sentence representations. Han Xiao [17] claims that the last layer is too close to the target function and is bi-

ased towards the pre-training task targets and hence, argues that the second-to-last layer is a better sentence representation. Other studies [10] and [13] showed that taking average of first and last layers leads to better sentence representations. In this study, we also aim to experiment with various pooling methods to contribute more data for the research community on which pooling method works best.

3 Solution

As DistilBERT proved to be promising in reducing memory footprint with minimal loss in language understanding, we decided to start off with an initial evaluation benchmark for BERT and DistilBERT. We evaluated these two models on Semantic Textual Similarity (STS) tasks using Spearman Correlation score as the main evaluation metric. Surprisingly, results showed that pretrained DistilBERT had consistently outperformed pretrained BERT on STS tasks. BERT scored an average of 53.73, while DistilBERT scored an average of 58.23. This shows DistilBERT to be a promising candidate as a pre-trained model to be used in our experiment.

We adopted a similar approach to unsupervised SimCSE [6] in applying contrastive learning to pre-trained language models. To perform contrastive learning in self-supervised fashion, we take sentences $\{x_i\}_{i=1}^m$ and use respective x_i as x_i^+ . During forward pass, dropouts from DistilBERT encoder is applied, resulting in embeddings that are similar but not identical for each pair of x_i and x_i^+ . We can denote these two embeddings as $h_i^z = f(x_i, z)$ and $h_i^{z'} = f(x_i, z')$ where z and z' are two different random dropout masks applied on x_i . The cosine similarity is calculated for positive pairs x_i and x_i^+ , and negative pairs x_i and x_j^+ . Finally, our training loss for contrastive learning is computed as:

$$\ell_i = -\log \frac{e^{\text{sim}(h_i, h_i^+)/\tau}}{\sum_{j=1}^N e^{\text{sim}(h_i, h_j^+)/\tau}}$$

where τ is a temperature hyperparameter that can be tuned for improving model performance. Here, minimizing the loss function will result in maximizing the similarity between positive pairs (increasing the numerator), and minimising the similarity between negative pairs (de-

creasing the denominator). The overall contrastive learning architecture using DistilBERT is shown below:

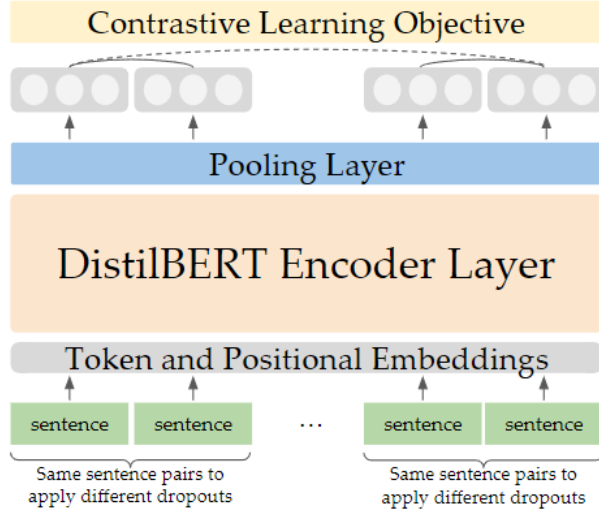


Figure 1: Overall DistilFACE Architecture. Similarity of final embeddings after the pooling layer are measured. Solid lines are positive examples, while dotted lines are negative examples.

4 Methodology

Pre-trained DistilBERT model with initial checkpoints from Transformers Hub is downloaded and used as the starting point for applying contrastive learning on DistilBERT.

DistilFACE model is trained using Wiki 1M sampled dataset obtained from SimCSE. With reference to common hyper-parameters used by previous works in BERT, DistilBERT, and SimCSE, we started with a default set of hyperparameters in this study shown in table below:

<u>Optimizer</u> Adam	<u>Maximal Length</u> 32	<u>Batch Size</u> 64	<u>Learning rate</u> 5e-05
<u>Pooling method</u> Average first last		<u>Temperature</u> 0.05	<u>Max Steps</u> 30,000

Figure 2: Default Hyperparameters

A grid search is performed to find the optimal set of hyperparameters, more explanation can be found in the Results & Analysis section. In grid search, we also perform evaluation to understand the effect of hyperparameters on STS task performance.

4.1 Success Metrics

In the evaluation phase, we used STS-12, STS-13, STS-14, STS-15 and STS-B as our evaluation set for measuring model performance on semantic tasks. Following the methodology in SimCSE, the SentEval engine by Conneau et. al. [3] was mainly used to evaluate model performance on these evaluation datasets, and Spearman Correlation is used as the evaluation metric. This metric helps us to determine the quality of the sentence embeddings produced for determining semantic similarity.

To ensure our Spearman correlation results are comparable with SimCSE, the final spearman correlation score for each dataset is calculated by first concatenating all subsets of data within the STS dataset, and then calculating Spearman correlation once, similar to the approach by SimCSE. This is slightly different from the default calculation from SentEval, which takes the average of Spearman correlation scores for each data subset within each STS dataset.

To measure success, after hyperparameter tuning is complete, the best hyperparameters will be used with max steps tuned to adjust for overfitting in the evaluation set. Spearman Correlation scores for the final DistilFACE will be compared with other models such as BERT, DistilBERT and SimCSE.

4.2 Further Enhancements

To make our model even more efficient and require less GPU memory, we also adopted the use of Automatic Mixed Precision (AMP). Instead of always storing weights and biases in 32-bit float format, 16 bit is sometimes used instead, such as when there are. This process of autocasting to lower precision where appropriate reduces wastage of GPU memory. An initial issue faced was that performance was noticeably worse. However, we realized that this was likely due to underflow, as the gradients can become too small to be stored via this format. This can be managed by applying a scaling factor to

gradients. The difference made by this scaling factor improved average scores across benchmarks by 5.6%. Overall, utilizing AMP resulted in a speed up of about 1.4x for training with a batch size of 64, to up to 3.2x for batch size of 256, and meant we needed a maximum of 8GB in GPU memory during our training process, enabling local training in line with our secondary project objectives.

5 Results & Analysis

We compare the evaluation results of our final DistilFACE model with other models:

Model	STS12	STS13	STS14	STS15	STSB	Avg.
BERT (Avg. first-last pooling)	39.70	59.37	49.68	66.03	53.87	53.73
DistilBERT (Avg. first-last pooling)	47.59	61.87	52.94	69.69	59.05	58.23
SimCSE-BERTbase	68.40	82.41	74.38	80.91	76.85	76.25
DistilFACE (Avg. last-4 pooling)	63.72	75.29	69.49	78.22	73.80	72.10

Figure 3: Comparison of Model Performances

From the above, we can clearly see that our application of contrastive learning to DistilBERT has been clearly successful based on our primary success metric of improved performance across STS tasks. DistilFACE was able to perform on average 34.2% and 23.8% higher in Spearman Correlation compared to BERT and DistilBERT respectively. We have also achieved our secondary goal of building a lightweight model, since DistilFACE, which is based on DistilBERT, is significantly smaller and faster than BERT (shown in figure below).

{Insert Table here}

We believe that our experiments were successful for two key reasons. The first is that our hypothesis in section 2.3 that the properties of contrastive learning are compatible with DistilBERT is logically sound. The second is effective hyperparameter tuning and experimentation with pooling methods enabled good performance. We detail the results and analysis of that process below.

5.1 Hyperparameter tuning

These are the hyperparameters that we tuned: Similarity Temperature, batch size, learning rate and pooling

method. Each hyperparameter is tested against their selected set of values, which are adapted from [6]. The optimal hyperparameters and evaluation score by training steps is shown below:

<u>Optimizer</u> Adam	<u>Maximal Length</u> 32	<u>Batch Size</u> 128	<u>Learning rate</u> 1e-05
<u>Pooling method</u> Average last 4		<u>Temperature</u> 0.05	<u>Max Steps</u> 20,000

Figure 4: Optimal Hyperparameters for Final DistilFACE




Figure 5: Spearman Corr. by Steps on STS datasets

5.2 Challenges Faced

A problem we encountered was when we sought to experiment with hyperparameter tuning frameworks like Ray and Optuna to enable more sophisticated methods. We did not anticipate issues as there are examples of successful applications to BERT models. However, these frameworks were found to be incompatible with our use case. Firstly, the trainer required labels from the dataset by default, which is not applicable for our unsupervised model. This caused the “compute_metric” function to be ignored by the tuning framework, which is not ideal given that we require the use of the SentEval toolkit for evaluation. Secondly, our custom model in evaluation mode also does not directly output evaluation loss either, instead such metrics are calculated separately outside the model via the

SentEval toolkit. This toolkit is also rather complex to be practically integrated within the requirements of the “compute_metric” function.

Hence, we instead adopt the grid-search methodology adopted by the authors in [6]. Each set of hyperparameter values are used in training for 30,000 steps, and then evaluated independently on these 5 datasets described above. We selected this number of steps as we found training loss to be stable by this threshold across the different hyperparameters (Appendix Figure 16 for training loss). This reasonable amount of steps also meant that using grid search proved sufficient, as gains from sophisticated methods like parallelisation via Ray and Optuna would be more limited.

5.3 Learning Rate

Learning rates of 1e-5, 3e-5 and 5e-5 were selected as hyperparameter values as adapted from [6]. Learning rate of 1e-5 is chosen as it resulted in the best average spearman correlation across the 5 datasets. As evident in Figure 16, the larger the learning rate, the greater the spearman correlation.

5.4 Similarity Temperature

Temperatures of 0.001, 0.01, 0.05, 0.1, 1 were tested, with 0.05 being chosen as it resulted in the best average spearman correlation across the 5 datasets, as shown in Figure 16. This parameter scales the inner product between two normed feature embeddings[11] when calculating the training objective. The temperature controls how sensitive the objective is to specific embedding locations [19], and as such this value needs to be neither too large nor small, and in this scenario optimally at 0.05.

5.5 Batch Size

Batch sizes of 64, 128, 256 were tested. Batch size of 128 is chosen as it resulted in the best average spearman correlation across the 5 datasets, as shown in Figure 16. However, as can also be seen from this figure, the model is not too sensitive to batch sizes provided a suitable learning rate is used, consistent with the findings from [6]. This provides further validation that contrastive learning need

not require very large batch sizes [6], unlike what was previously thought[2]. This also further supports our overarching hypothesis that contrastive learning is suitable to be applied to more light weight models.

5.6 Pooling Methods

We experimented with various pooling methods taking reduce functions: average and max, as well as the selection of hidden layers: last hidden, second-to-last, all hidden, first and last, last two, last four and concatenation of last four layers. Our findings are as follows:

Best Pooling Methods. Concat and average of the last four layers were the best two performing pooling methods. This is consistent with BERT [12] where concat of the last four layers worked best for them.

Next Best Pooling Methods. Authors [10] [13] mentioned that taking the average of first and last leads to better results. We find this consistent with our results with average of first and last to be one of the better performing pooling methods out of all we have experimented.

Never too close. Our results show that the last hidden is the third best performing pooling method, better than the second-to-last layer. This discredits the claim from Han Xiao [17] that the last layer is too close to pre-trained tasks and the second-to-last layer is a better sentence representation. Here, we find this to be false.

6 Conclusion

Overall, we believe our project to be a success using the success metrics in section 4.1, namely Spearman Correlation on STS tasks. This can be seen from the results and rationale in section 5. Hence we hope our DistilFACE model can help lay the groundwork for future work involving building lightweight models for such tasks using contrastive and knowledge distillation.

References

- [1] H. G. Abreha, M. Hayajneh, and M. A. Serhani. Federated learning in edge computing: A systematic survey. *Sensors*, 22(2):450, Jan 2022. url: <http://dx.doi.org/10.3390/s22020450>.

- [2] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton. A simple framework for contrastive learning of visual representations, 2020, [arXiv:2002.05709](#).
- [3] A. Conneau and D. Kiela. Senteval: An evaluation toolkit for universal sentence representations, 2018, [arXiv:1803.05449](#).
- [4] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2018, [arXiv:1810.04805](#).
- [5] M. Farahmand. Pre-trained word embeddings or embedding layer?, 2019. url: [link](#).
- [6] T. Gao, X. Yao, and D. Chen. Simcse: Simple contrastive learning of sentence embeddings, 2021, [arXiv:2104.08821](#).
- [7] R. D. Hjelm, A. Fedorov, S. Lavoie-Marchildon, K. Grewal, P. Bachman, A. Trischler, and Y. Bengio. Learning deep representations by mutual information estimation and maximization, 2018, [arXiv:1808.06670](#).
- [8] A. Kashyap, L. Han, J. Sleeman, T. Satyapanich, S. Gandhi, and T. Finin. Robust semantic text similarity using lsa, machine learning, and linguistic resources. *Lang Resources & Evaluation*, 50, 2016. url: [link](#).
- [9] Z. Kira. Cs7643 lecture 18: Unsupervised and self-supervised learning, 2022. url: [Lecture 18](#).
- [10] B. Li, H. Zhou, J. He, M. Wang, Y. Yang, and L. Li. On the sentence embeddings from pre-trained language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9119–9130, Online, Nov. 2020. Association for Computational Linguistics. url: <https://aclanthology.org/2020.emnlp-main.733>.
- [11] S. Li, J. Xu, and B. Hooi. Probabilistic contrastive loss for self-supervised learning, 2021, [arXiv:2112.01642](#).
- [12] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov. Roberta: A robustly optimized bert pretraining approach, 2019, [arXiv:1907.11692](#).
- [13] N. Reimers and I. Gurevych. Sentence-bert: Sentence embeddings using siamese bert-networks, 2019, [arXiv:1908.10084](#).
- [14] V. Sanh, L. Debut, J. Chaumond, and T. Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter, 2019, [arXiv:1910.01108](#).
- [15] S. Sun, Z. Gan, Y. Cheng, Y. Fang, S. Wang, and J. Liu. Contrastive distillation on intermediate representations for language model compression, 2020, [arXiv:2009.14167](#).
- [16] M. S. Tony Peng. The staggering cost of training sota ai models, 2019. url: [link](#).
- [17] H. Xiao. Why not the last hidden layer? why second-to-last?, 2018. url: [Why not the last hidden layer](#).
- [18] Y. Yan, R. Li, S. Wang, F. Zhang, W. Wu, and W. Xu. Consert: A contrastive framework for self-supervised sentence representation transfer, 2021, [arXiv:2105.11741](#).
- [19] O. Zhang, M. Wu, J. Bayrooti, and N. Goodman. Temperature as uncertainty in contrastive learning, 2021, [arXiv:2110.04403](#).